

Visión por Computador

CNN

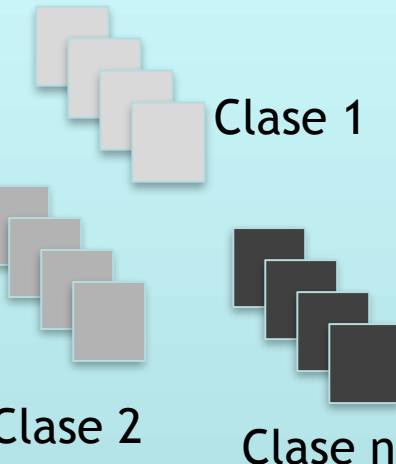
PhD. Carlos Fernando Montoya Cubas



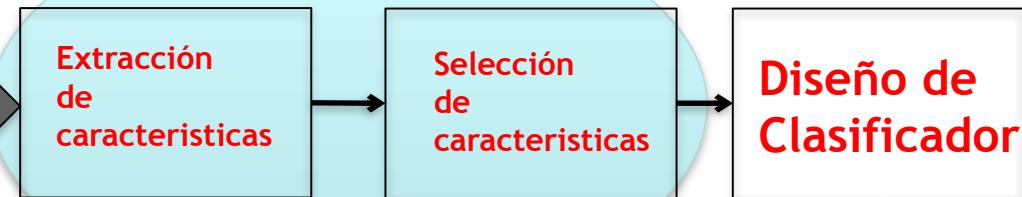
Motivación

PhD. Carlos Fernando Montoya Cubas

Imágenes de diferentes clases conocidas



Representación



Etapa 1: CAPACITACIÓN

Etapa 2: PRUEBAS

Aporte : Prueba de Imagen



MOTIVACIÓN

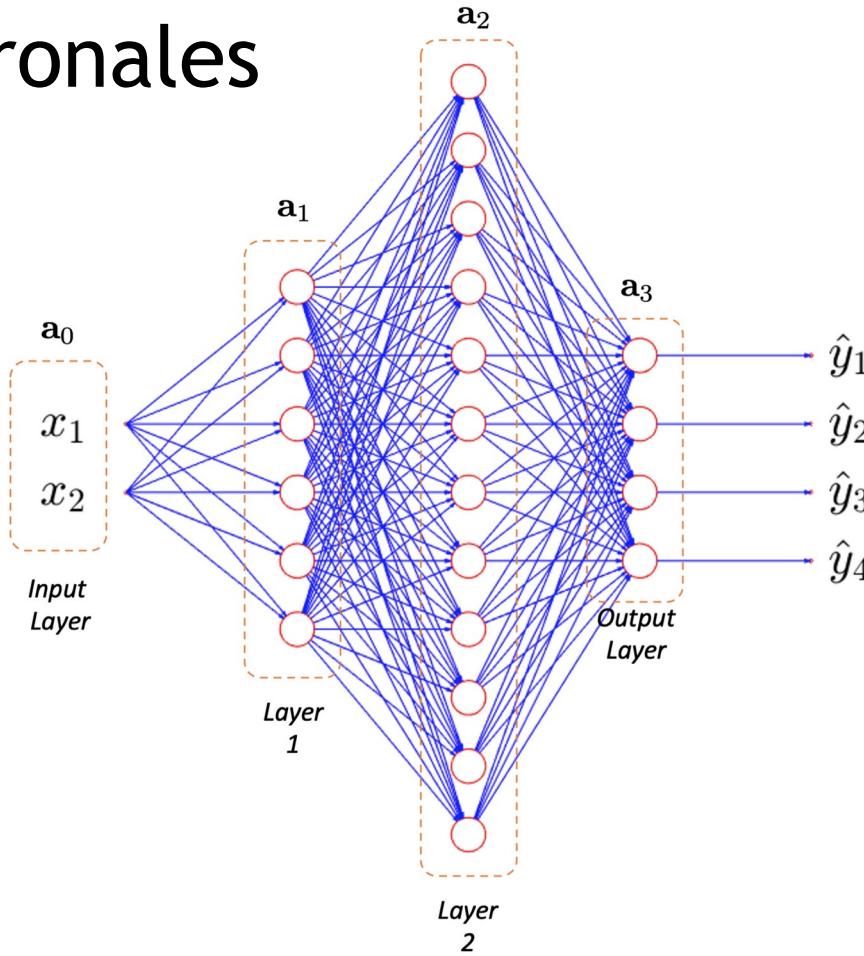
- Características manuales vs. aprendidas
- LBP, Haralick, Fourier, Gabor, etc

or

- Representación aprendida a partir de datos de entrenamiento

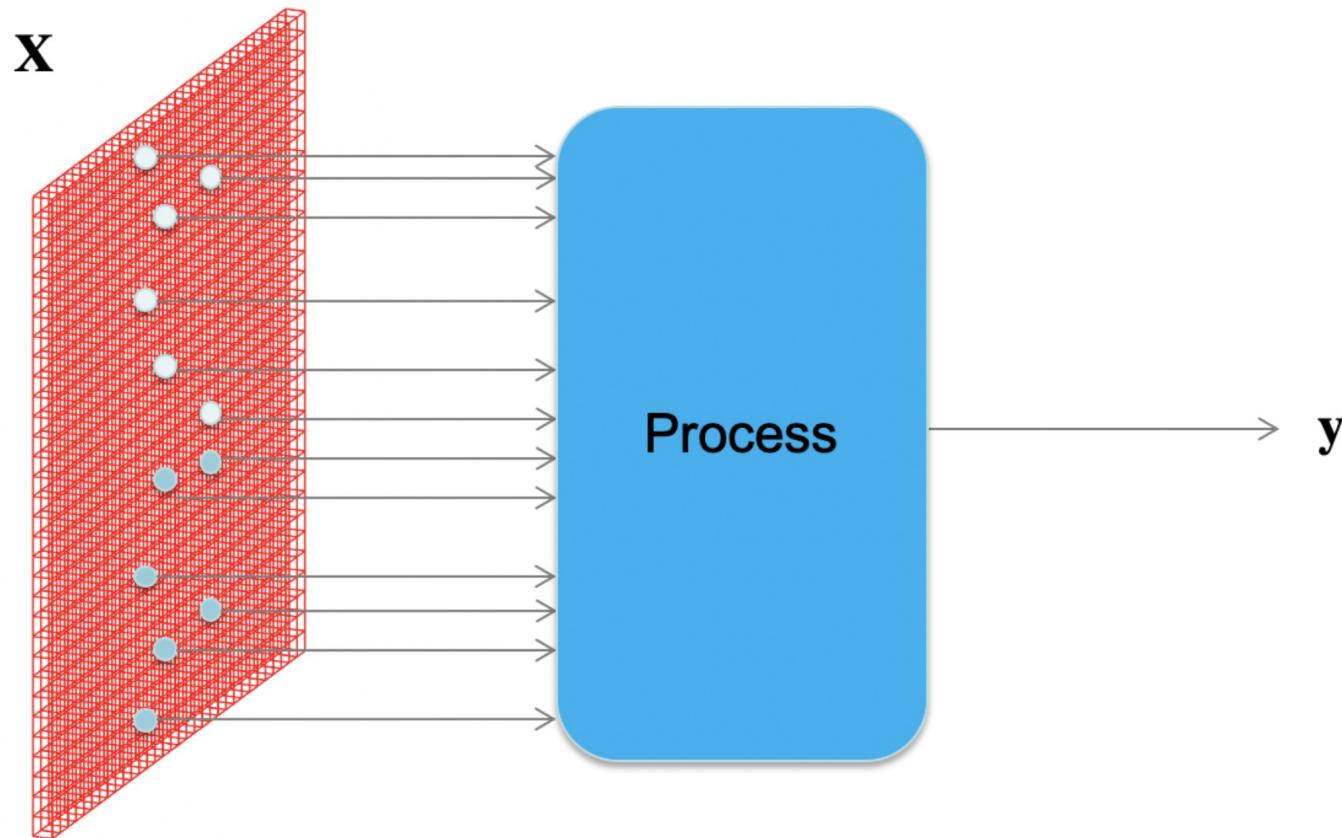
MOTIVACIÓN

- Redes neuronales



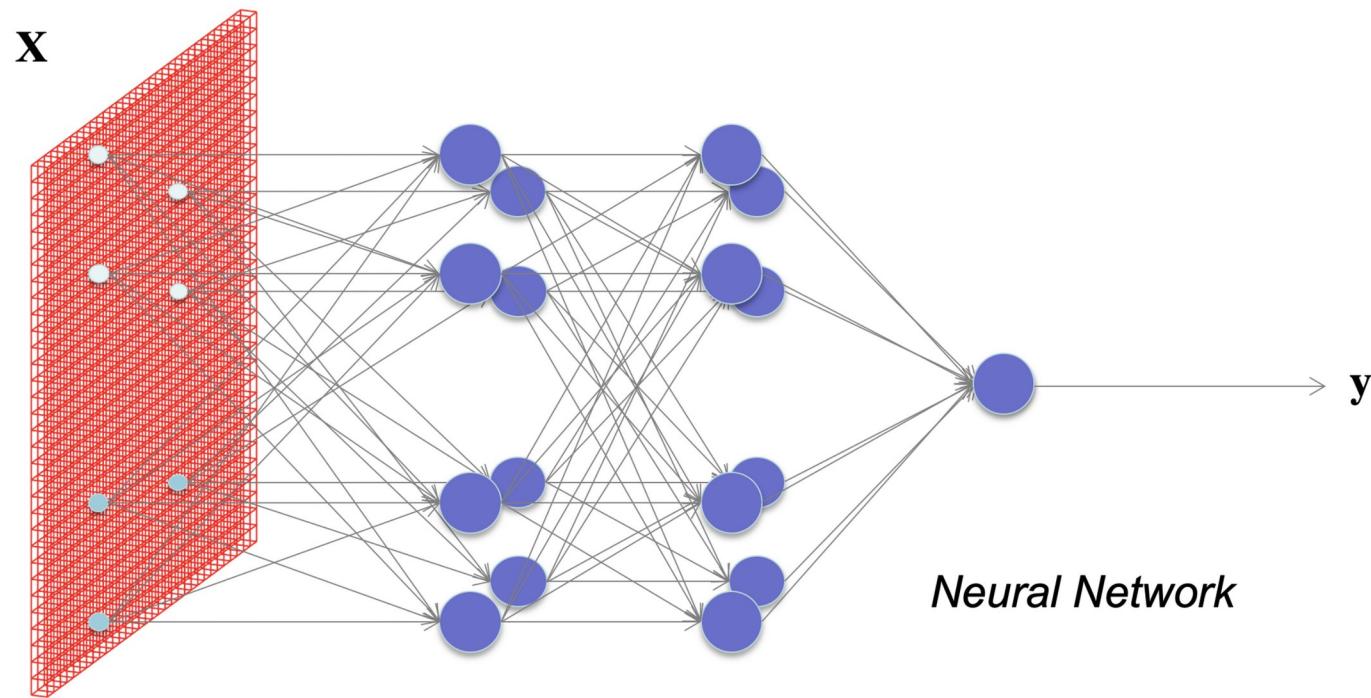
MOTIVACIÓN

- ¿Cómo procesar una imagen con redes neuronales?



MOTIVACIÓN

- ¿Qué sucede si cada píxel es una entrada?



¿Cuántas neuronas y pesos?

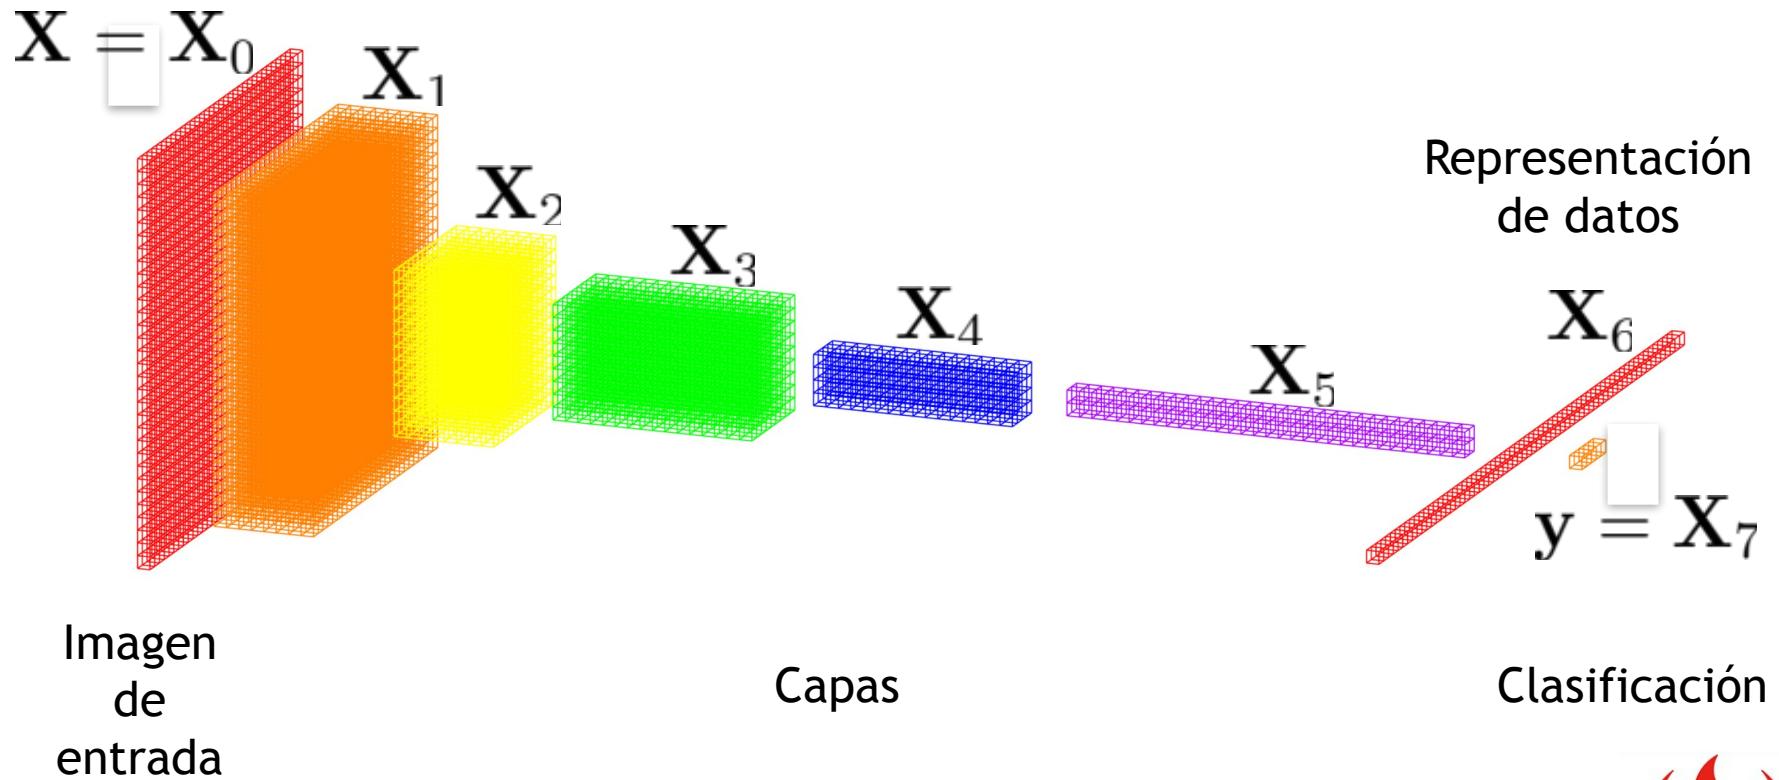


Básicos



DEEP LEARNING BASICS (CNN)

Redes neuronales convolucionales



PhD. Carlos Fernando Montoya Cubas

DEEP LEARNING BASICS (CNN)

Redes neuronales convolucionales

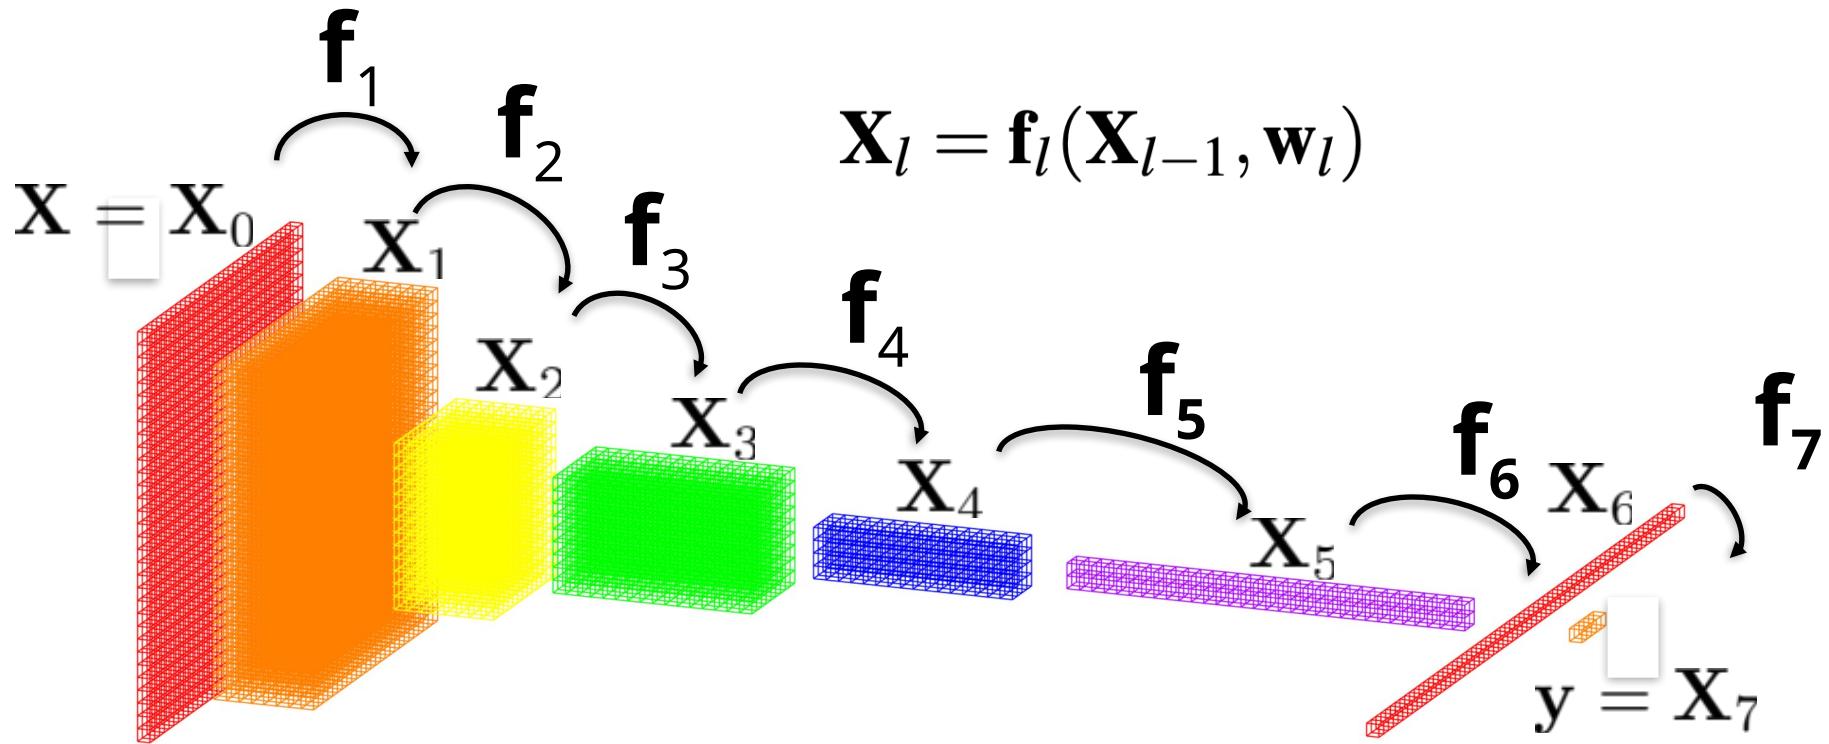


Imagen
de
entrada

Capas

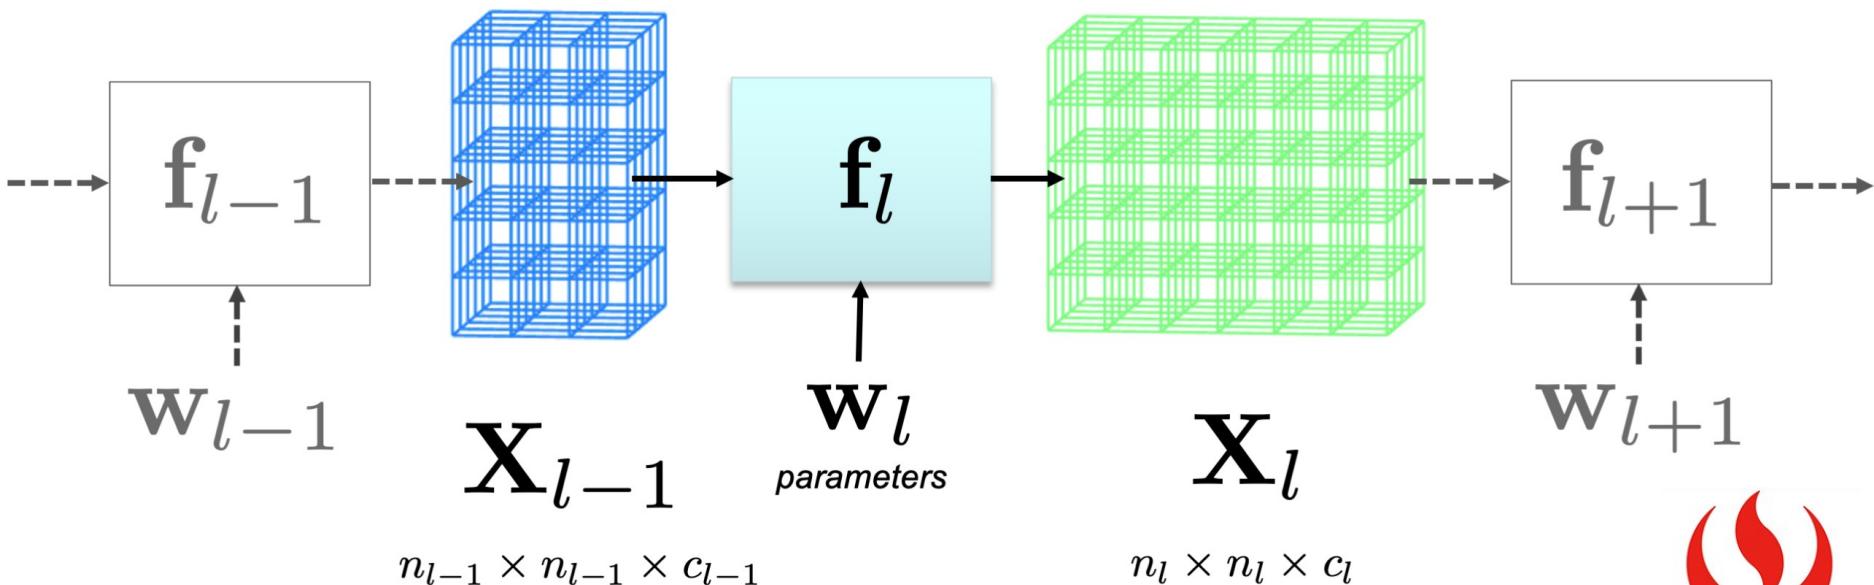
Clasificación

PhD. Carlos Fernando Montoya Cubas

DEEP LEARNING BASICS

La red neuronal convolucional, puede entenderse como un conjunto de capas, en el que cada capa procesa una imagen de entrada \mathbf{X} para obtener una salida imagen \mathbf{Y} :

$$\mathbf{X}_l = \mathbf{f}_l(\mathbf{X}_{l-1}, \mathbf{w}_l)$$



DEEP LEARNING: Capas típicas

- Capa de convolución
- Capa de agrupación
- Unidad Lineal Rectificada
- Totalmente conectado y SoftMax



DEEP LEARNING: Capas típicas

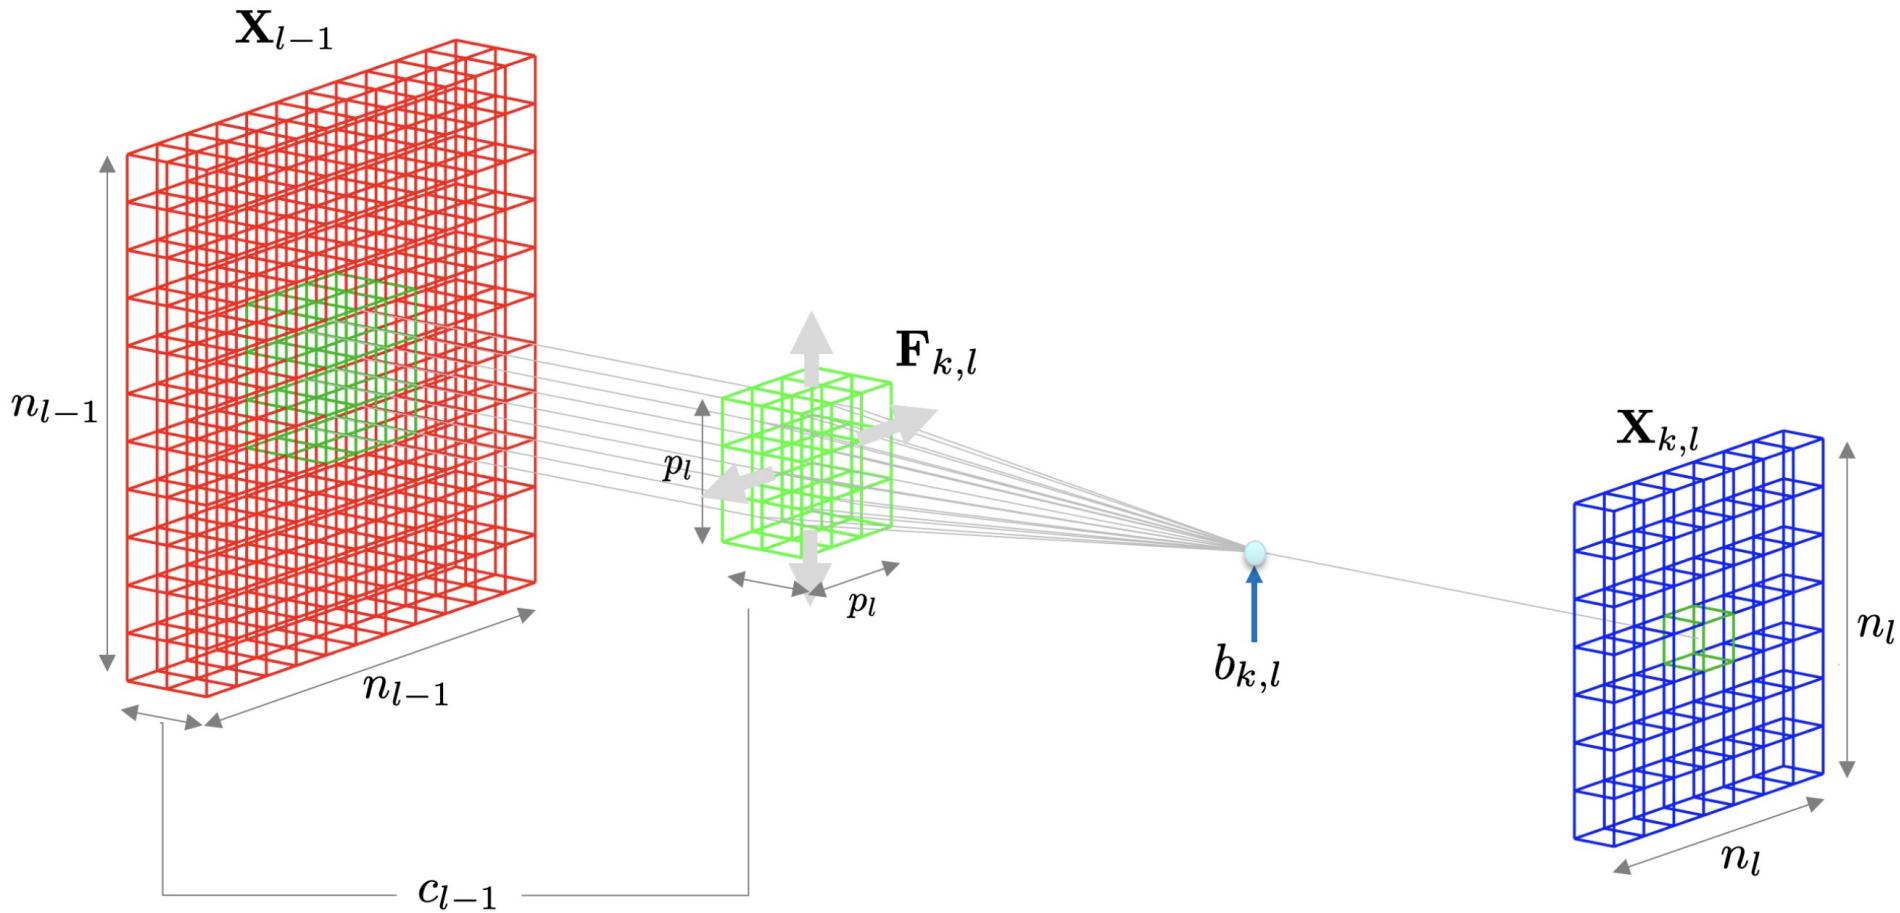
- Capa de convolución
- Capa de pooling
- Unidad Lineal Rectificada
- Totalmente conectado y SoftMax

DEEP LEARNING: Capa de convolución

$$\mathbf{X}_{k,l} = \mathbf{X}_{l-1} * \mathbf{F}_{k,l} + b_{k,l} \quad \text{for } k = 1 \dots m_l$$

$$X_{k,l}(i, j) = b_{k,l} + \sum_{u=1}^{p_l} \sum_{v=1}^{p_l} \sum_{w=1}^{q_l} X_{w,l-1}(i+u, j+v) F_{k,l}(u, v, w)$$

DEEP LEARNING: Capa de convolución



PhD. Carlos Fernando Montoya Cubas

DEEP LEARNING: Capas típicas

- Capa de convolución
- Capa de agrupación
- Unidad Lineal Rectificada
- Totalmente conectado y SoftMax

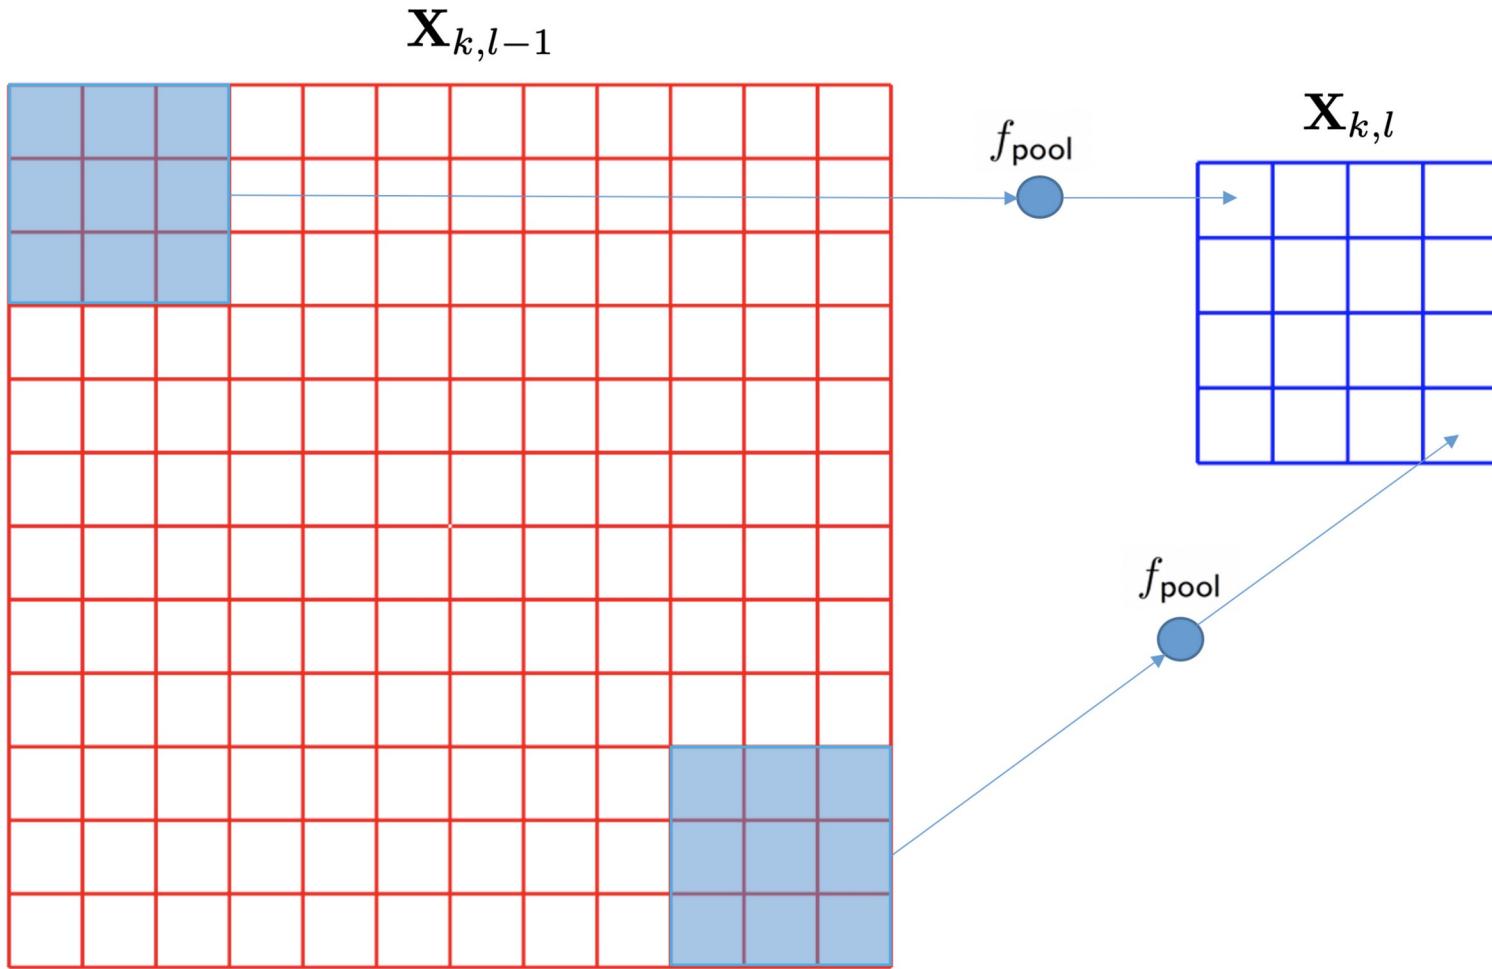


DEEP LEARNING: Capa de pooling

$$X_{k,l}(i,j) = f_{\text{pool}}\{X_{k,l-1}(u,v) : (u,v) \in \Omega(i,j)\}$$



DEEP LEARNING: Capa de pooling



DEEP LEARNING: Capas típicas

- Capa de convolución
- Capa de pooling
- Unidad Lineal Rectificada
- Totalmente conectado y SoftMax

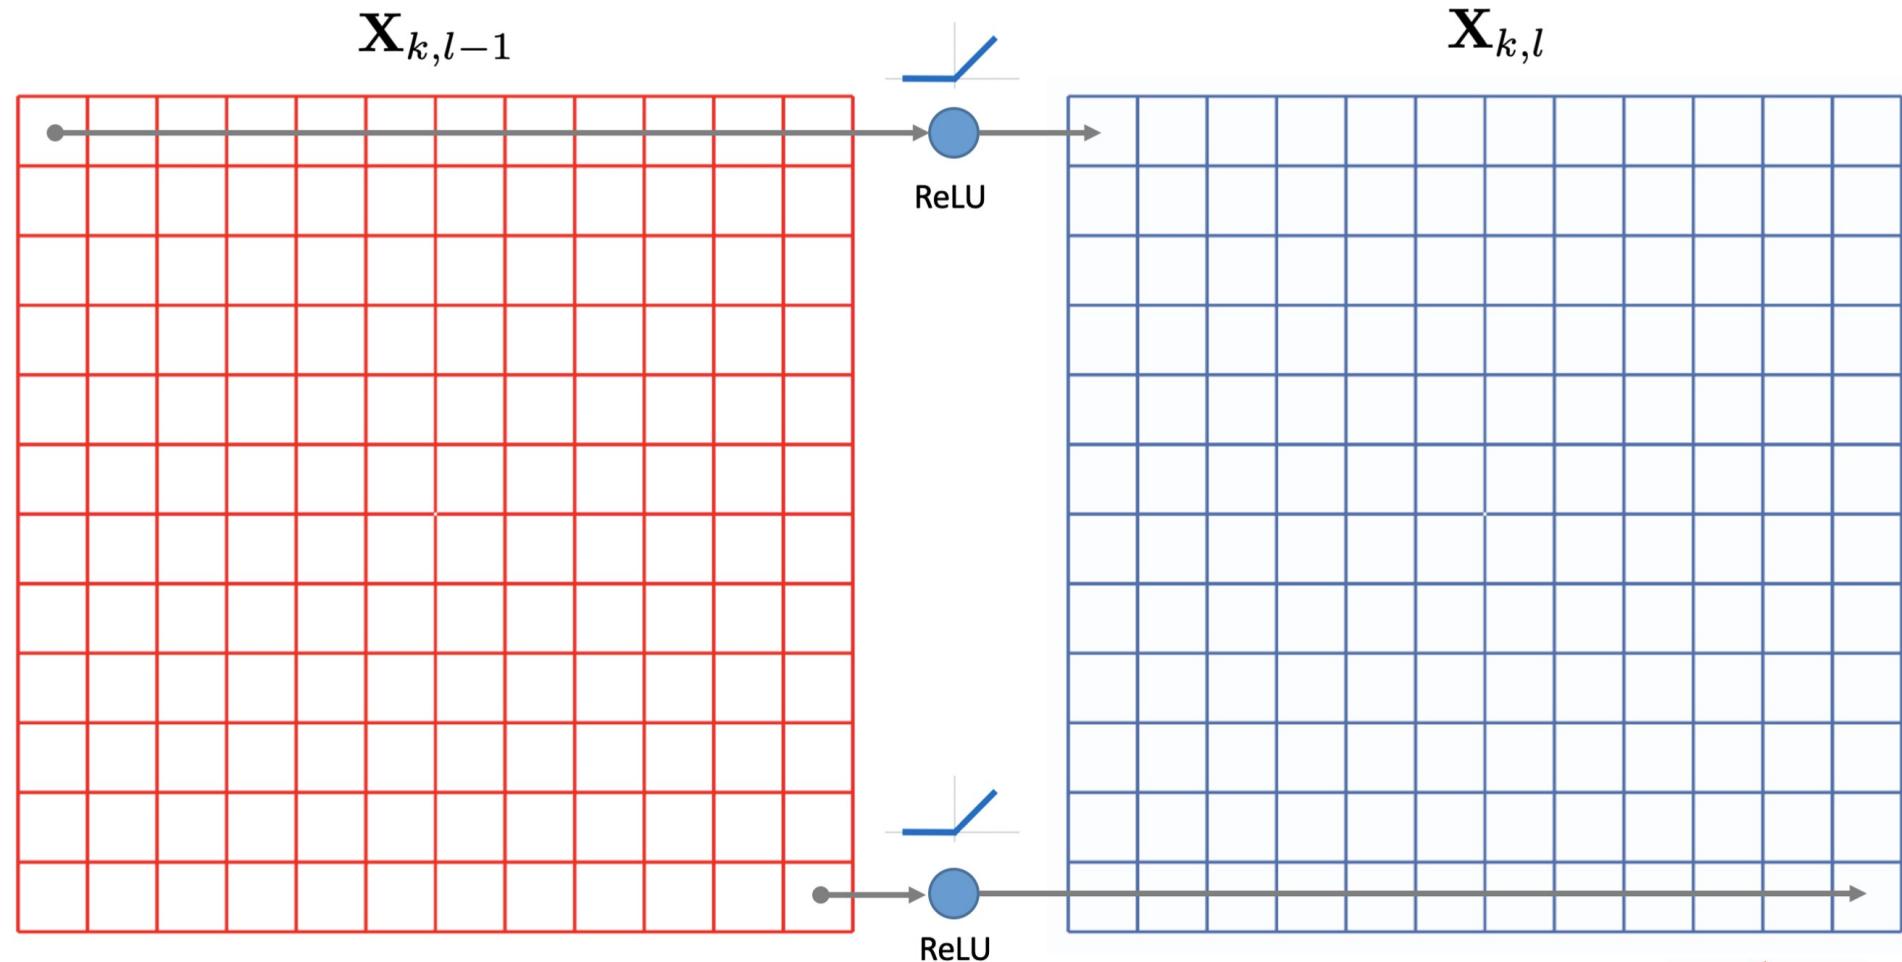


DEEP LEARNING: Unidad Lineal Rectificada(ReLU)

$$X_{k,l}(i,j) = \max\{0, X_{k,l-1}(i,j)\}$$



DEEP LEARNING: Unidad Lineal Rectificada (ReLU)

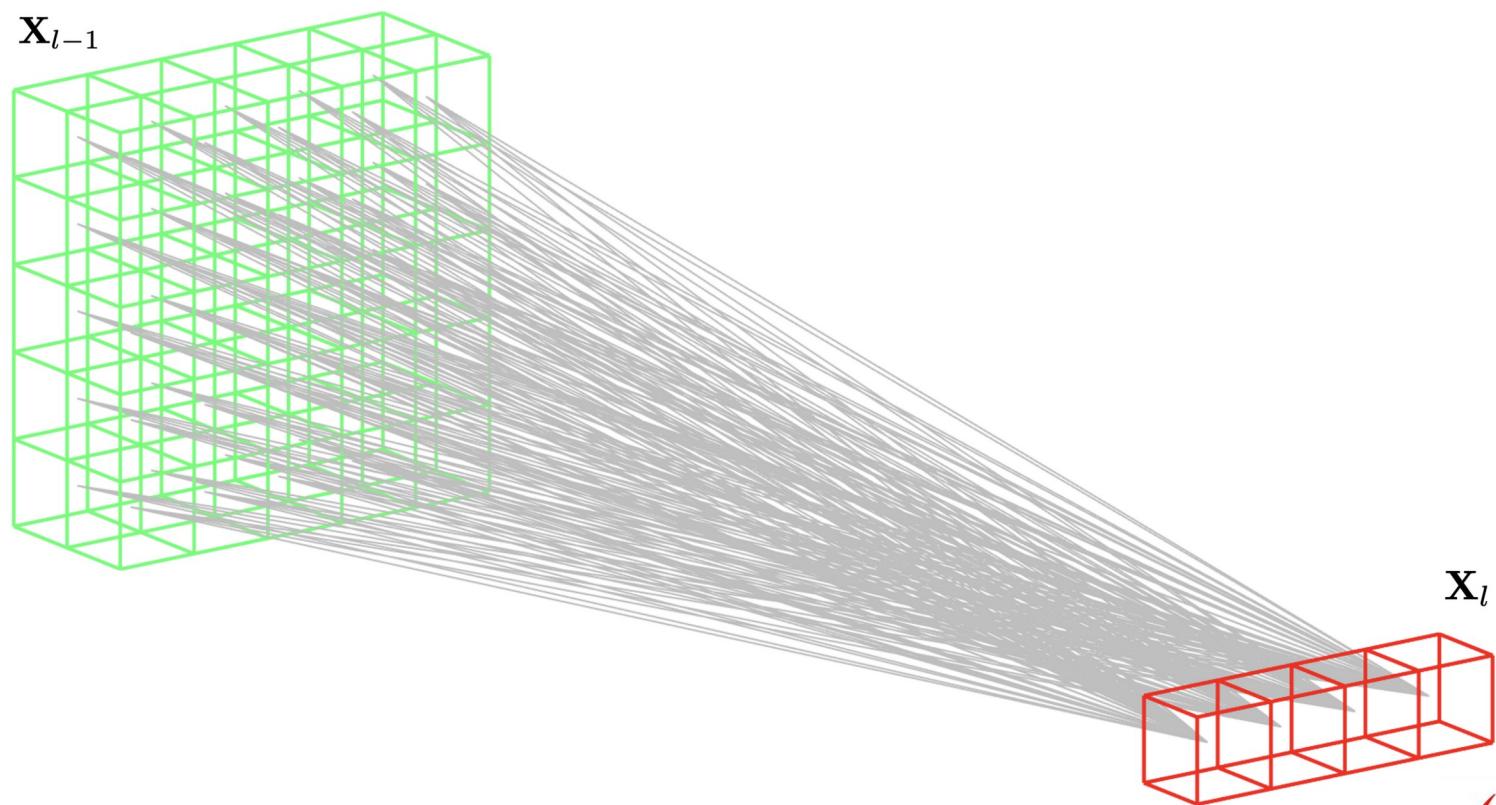


PhD. Carlos Fernando Montoya Cubas

DEEP LEARNING: Capas típicas

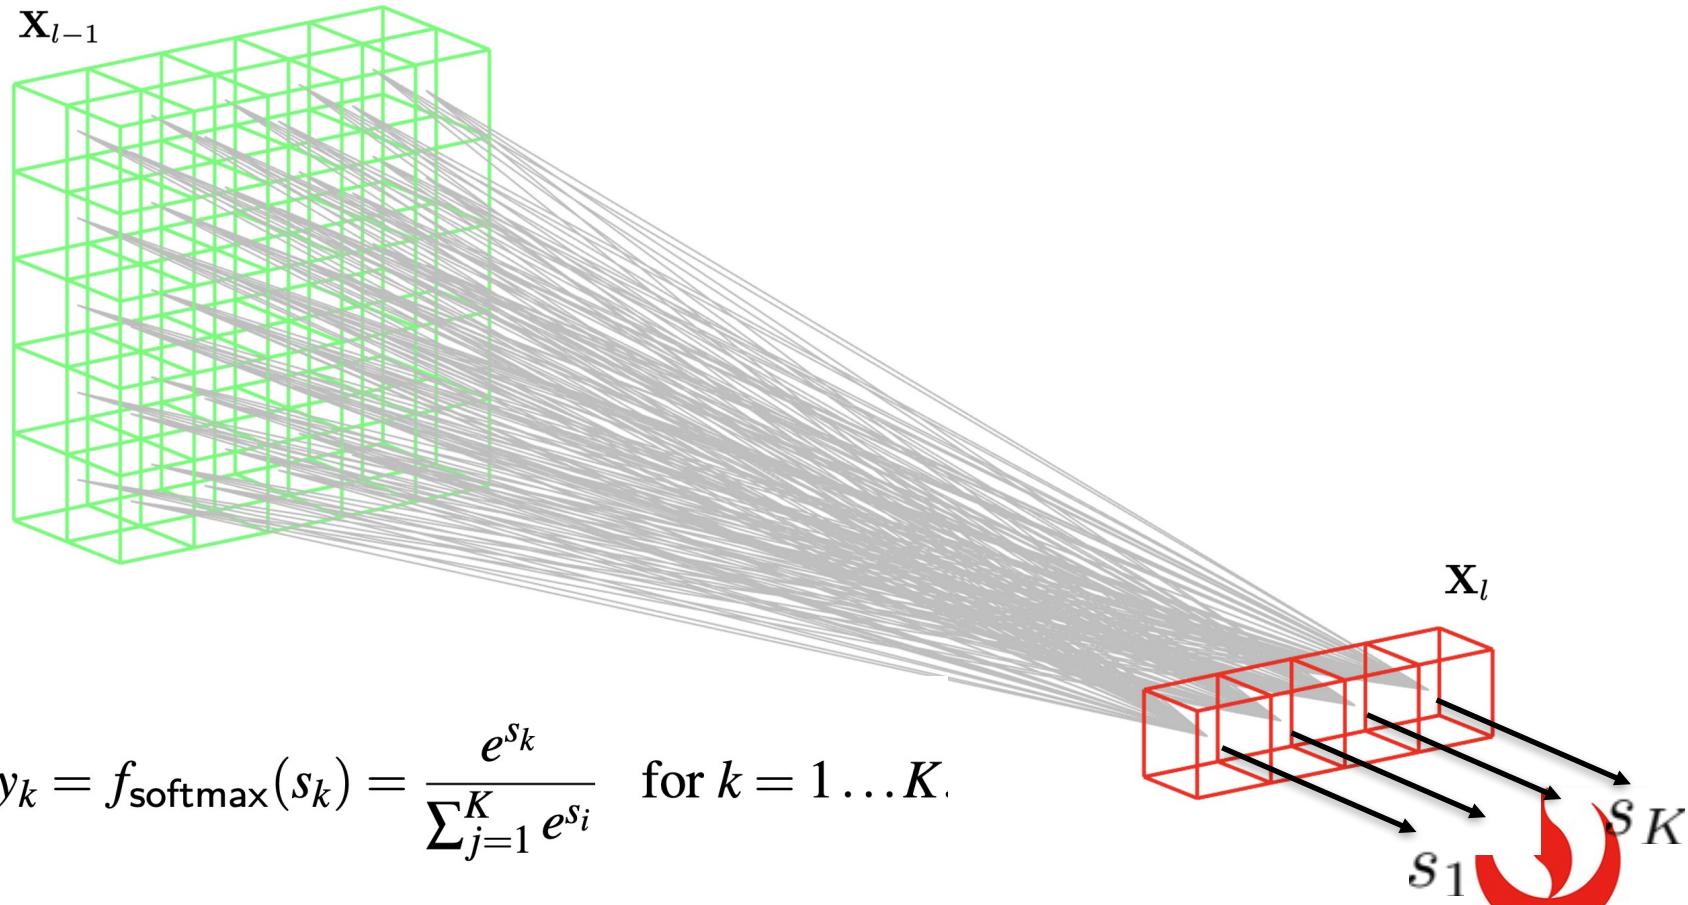
- Capa de convolución
- Capa de pooling
- Unidad Lineal Rectificada
- Totalmente conectado y SoftMax

DEEP LEARNING: Completamente conectado



PhD. Carlos Fernando Montoya Cubas

DEEP LEARNING: Completamente conectado + SoftMax



PhD. Carlos Fernando Montoya Cubas

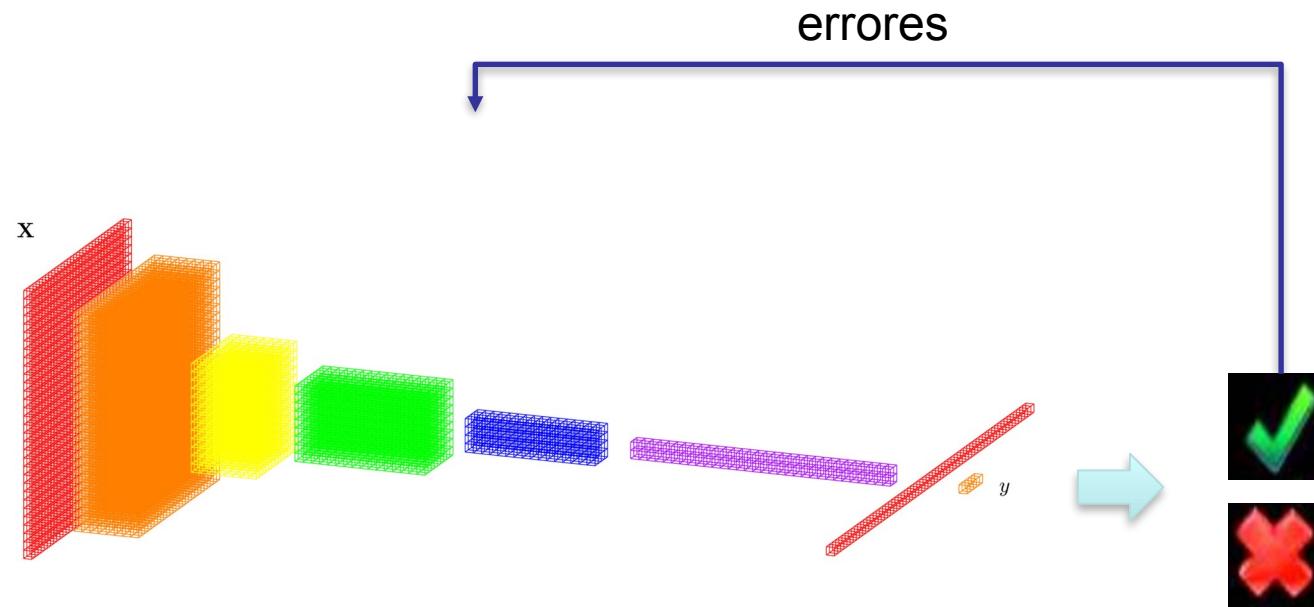
Training

DEEP LEARNING: Training

perros

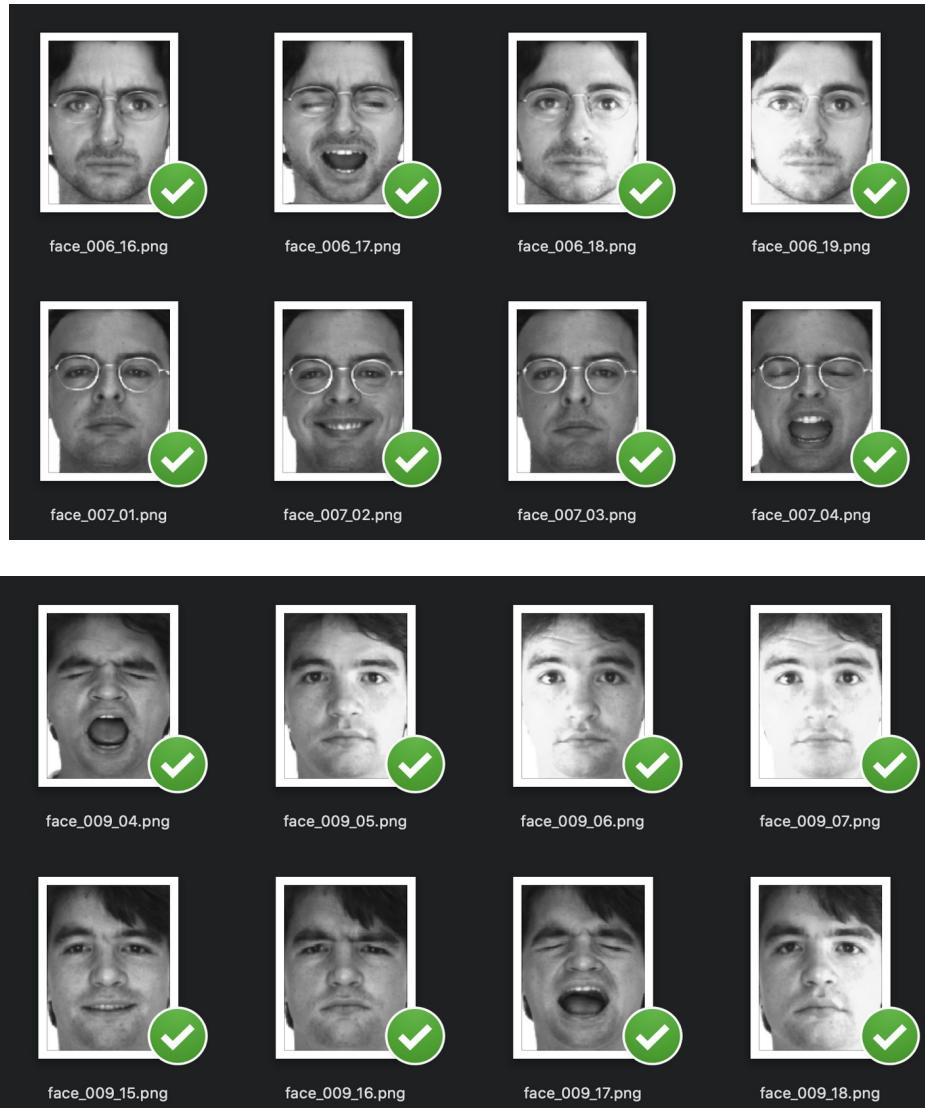


gatos

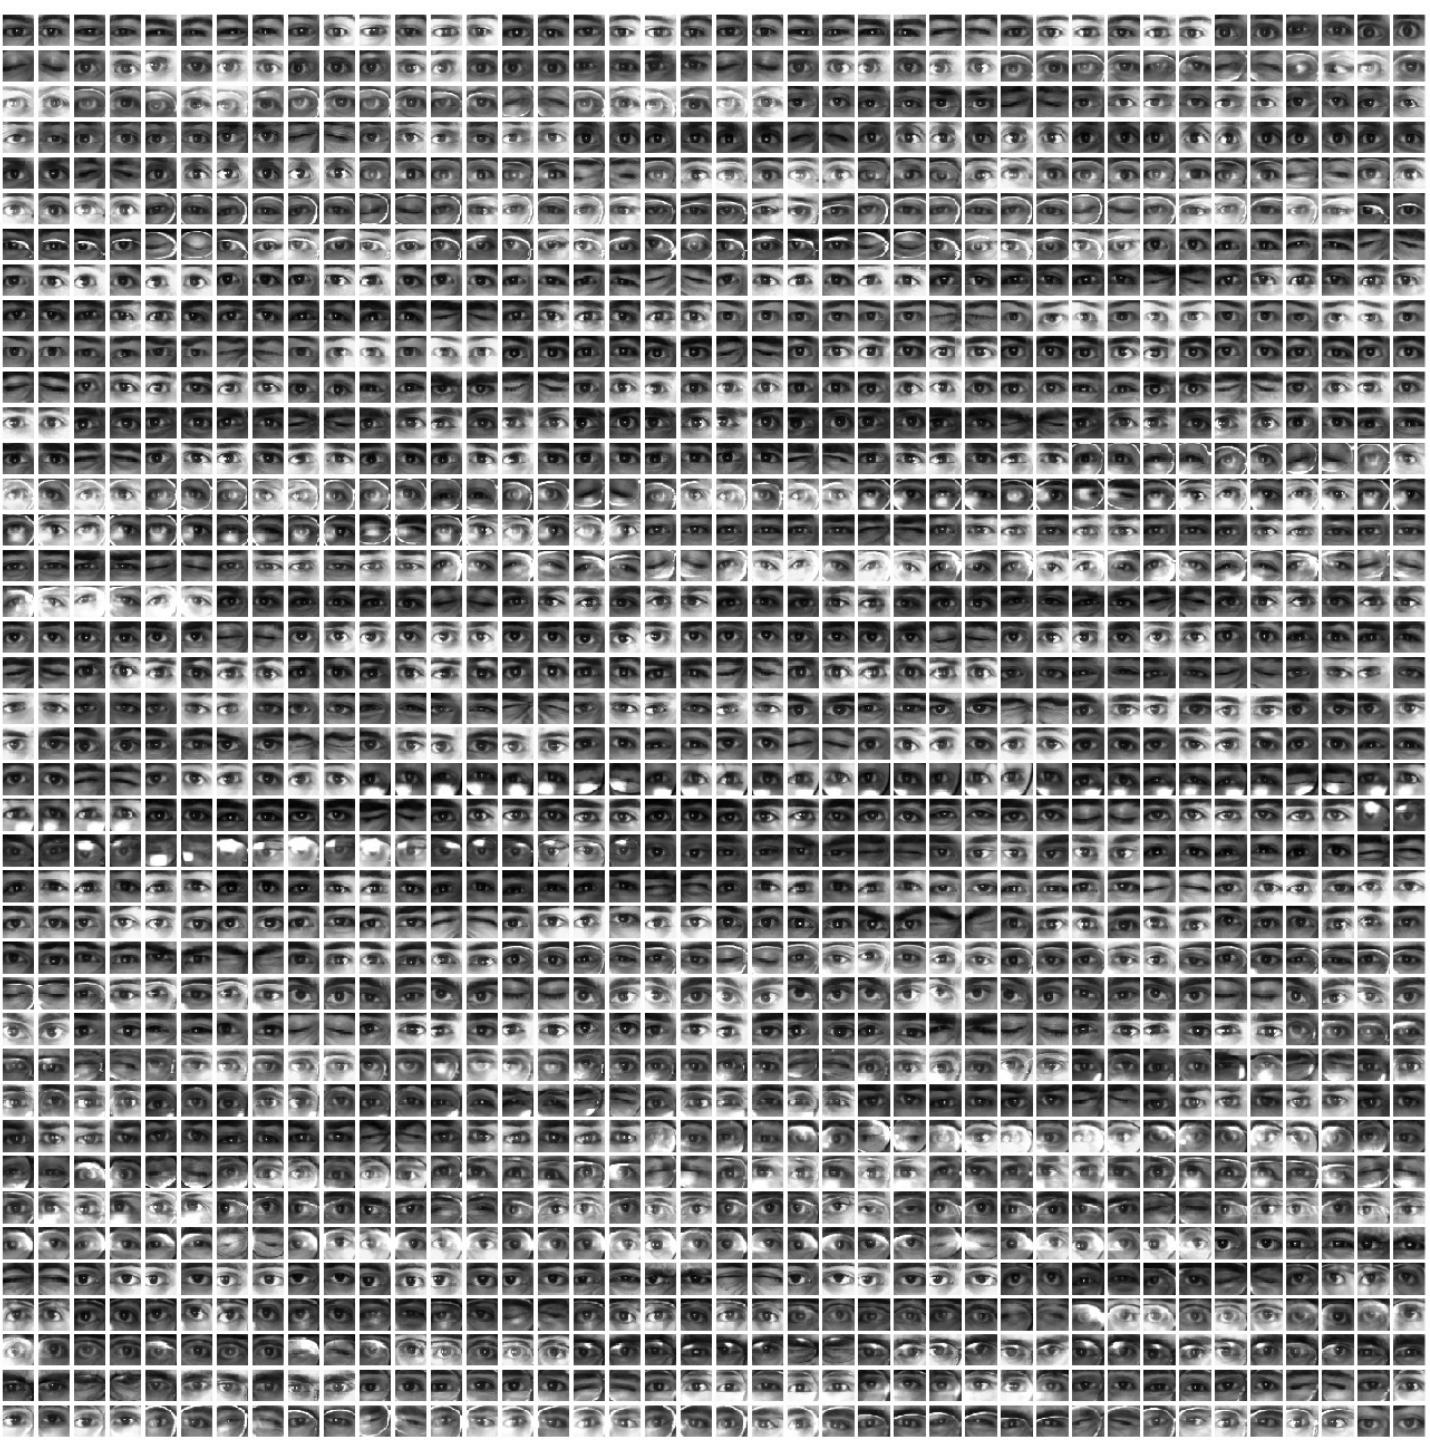


Ejemplo

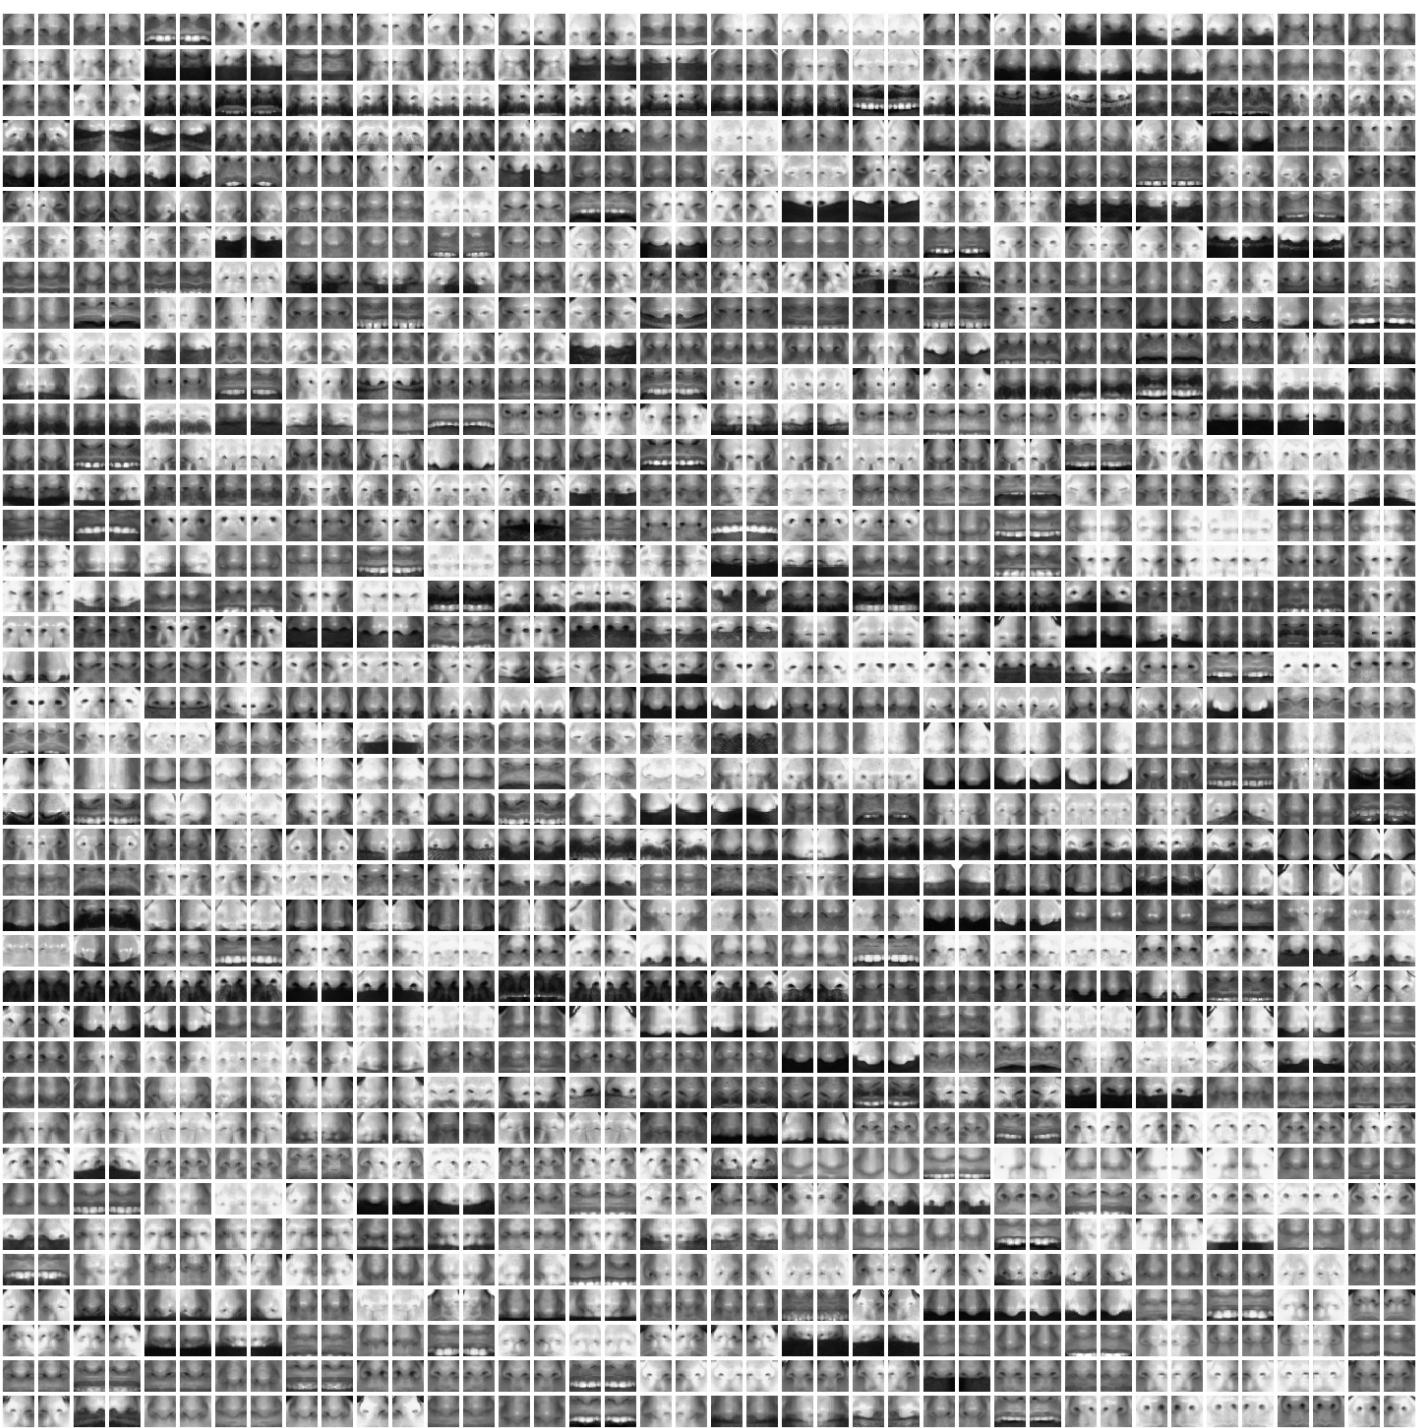
Reconocedor de ojos y nariz



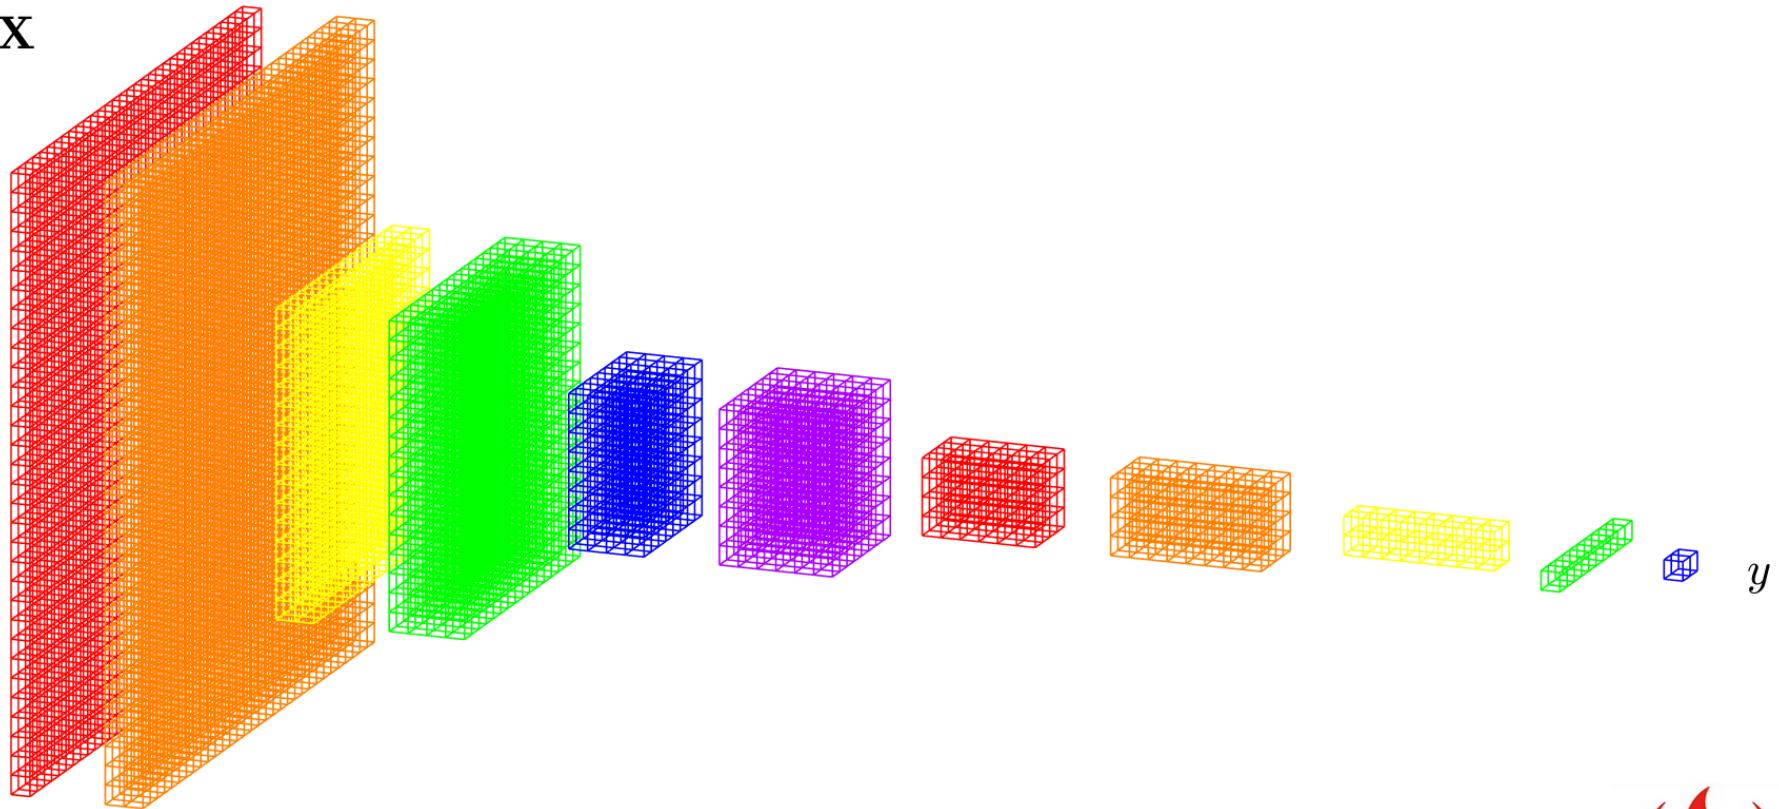
Ojos



Nariz



X



Implementation in PyTorch

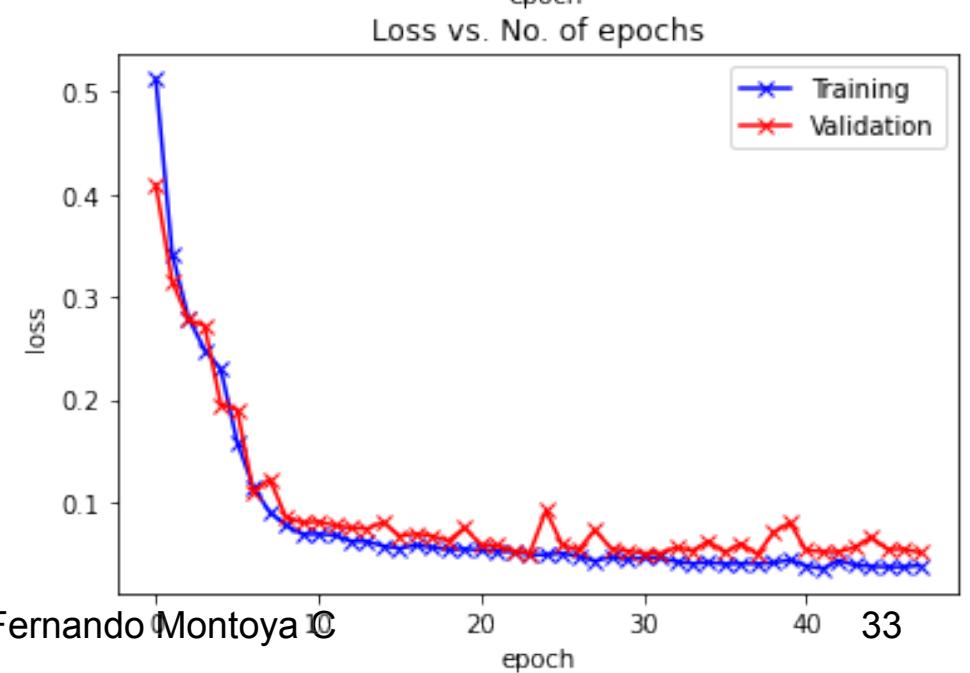
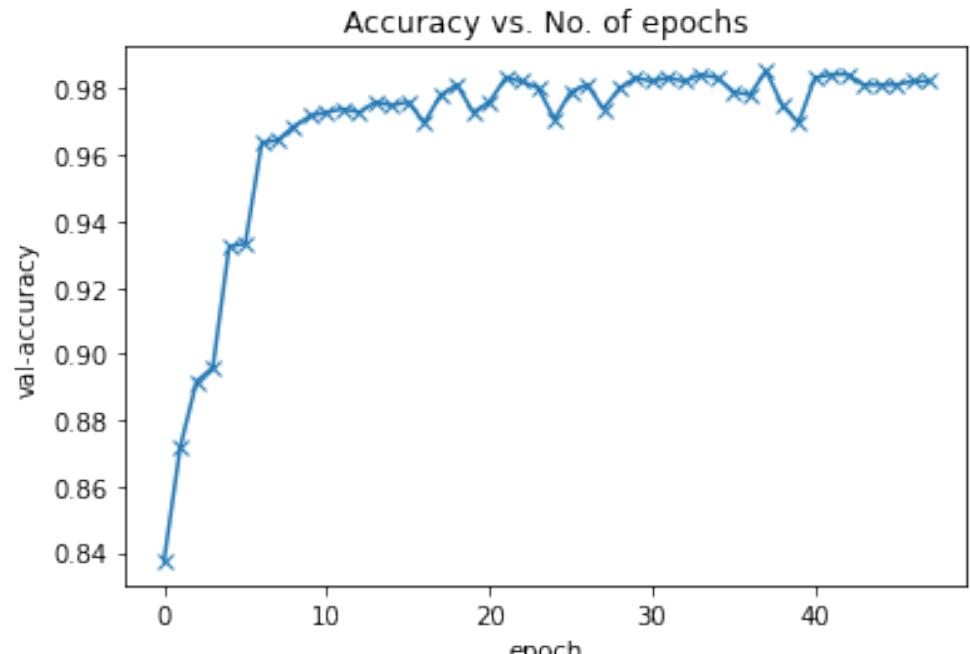
```
class CNN_Classification(ImageClassificationBase):
    def __init__(self):
        super().__init__()
        self.network = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=4, kernel_size = 5, stride = 1, padding = 0),
            nn.ReLU(),
            nn.Conv2d(in_channels=4, out_channels=8, kernel_size = 5, stride = 6, padding = 0),
            nn.ReLU(),
            nn.MaxPool2d(2,2),
            nn.Flatten(),
            nn.Linear(32,2)
        )
CNN_Classification(
    network): Sequential(
    (0): Conv2d(3, 4, kernel_size=(5, 5), stride=(1, 1))
    (1): ReLU()
    (2): Conv2d(4, 8, kernel_size=(5, 5), stride=(6, 6))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Flatten(start_dim=1, end_dim=-1)
    (6): Linear(in_features=32, out_features=2, bias=True)
)
```

Epoch	Train-Loss	Val-Loss	Val-Acc	Best	Time [sec]
0	0.5126	0.4093	0.8375	***	1.8
1	0.3427	0.3156	0.8719	***	1.8
2	0.2797	0.2781	0.8917	***	1.8
3	0.2486	0.2719	0.8958	***	1.8
4	0.2295	0.1950	0.9323	***	1.8
5	0.1588	0.1904	0.9333	***	1.8
6	0.1145	0.1096	0.9635	***	1.8
7	0.0904	0.1229	0.9646	***	1.8
8	0.0776	0.0855	0.9687	***	1.8
9	0.0688	0.0801	0.9719	***	1.8
10	0.0689	0.0811	0.9729	***	1.8
11	0.0681	0.0774	0.9740	***	1.8
12	0.0620	0.0750	0.9729		1.8
13	0.0617	0.0742	0.9760	***	1.8
14	0.0572	0.0806	0.9750		1.8
15	0.0551	0.0670	0.9760		1.8
16	0.0580	0.0692	0.9698		1.8
17	0.0560	0.0671	0.9781	***	1.8
18	0.0538	0.0624	0.9813	***	1.8
19	0.0544	0.0752	0.9729		1.8
20	0.0533	0.0582	0.9760		1.8
21	0.0522	0.0578	0.9833	***	1.9
22	0.0519	0.0526	0.9823		1.8
23	0.0489	0.0503	0.9802		1.8
24	0.0488	0.0918	0.9708		1.8
25	0.0504	0.0582	0.9792		1.8
26	0.0470	0.0535	0.9813		1.8
27	0.0428	0.0728	0.9740		1.8
28	0.0463	0.0551	0.9802		1.8
29	0.0443	0.0527	0.9833	***	1.8
30	0.0459	0.0504	0.9823		1.8
31	0.0461	0.0501	0.9833		1.8
32	0.0420	0.0557	0.9823		1.8
33	0.0402	0.0528	0.9844	***	1.8
34	0.0414	0.0619	0.9833		1.8
35	0.0397	0.0510	0.9792		1.9
36	0.0403	0.0597	0.9781		1.8
37	0.0400	0.0491	0.9854	***	1.8
38	0.0410	0.0704	0.9750		1.8
39	0.0438	0.0802	0.9698		1.8
40	0.0382	0.0532	0.9833		1.9
41	0.0350	0.0517	0.9844		1.8
42	0.0432	0.0516	0.9844		1.8
43	0.0387	0.0561	0.9813		1.8
44	0.0374	0.0660	0.9813		1.8
45	0.0370	0.0540	0.9813		1.8
46	0.0369	0.0544	0.9823		1.8
47	0.0383	0.0509	0.9823		1.8

*** Early stop after 10 epochs with no improvement

Best model saved best_model.pt (val_acc = 0.9854 in epoch = 37)
 Last model saved last_model.pt (val_acc = 0.9823 in epoch = 47)

Training Time: 85.82 sec



Epoch	Train-Loss	Val-Loss	Val-Acc	Best	Time [sec]
0	0.5126	0.4093	0.8375	***	1.8
1	0.3427	0.3156	0.8719	***	1.8
2	0.2797	0.2781	0.8917	***	1.8
3	0.2486	0.2719	0.8958	***	1.8
4	0.2295	0.1950	0.9323	***	1.8
5	0.1588	0.1904	0.9333	***	1.8
6	0.1145	0.1096	0.9635	***	1.8
7	0.0904	0.1229	0.9646	***	1.8
8	0.0776	0.0855	0.9687	***	1.8
9	0.0688	0.0801	0.9719	***	1.8
10	0.0689	0.0811	0.9729	***	1.8
11	0.0681	0.0774	0.9740	***	1.8
12	0.0620	0.0750	0.9729		1.8
13	0.0617	0.0742	0.9760	***	1.8
14	0.0572	0.0806	0.9750		1.8
15	0.0551	0.0670	0.9760		1.8
16	0.0580	0.0692	0.9698		1.8
17	0.0560	0.0671	0.9781	***	1.8
18	0.0538	0.0624	0.9813	***	1.8
19	0.0544	0.0752	0.9729		1.8
20	0.0533	0.0582	0.9760		1.8
21	0.0522	0.0578	0.9833	***	1.9
22	0.0519	0.0526	0.9823		1.8
23	0.0489	0.0503	0.9802		1.8
24	0.0488	0.0918	0.9708		1.8
25	0.0504	0.0582	0.9792		1.8
26	0.0470	0.0535	0.9813		1.8
27	0.0428	0.0728	0.9740		1.8
28	0.0463	0.0551	0.9802		1.8
29	0.0443	0.0527	0.9833	***	1.8
30	0.0459	0.0504	0.9823		1.8
31	0.0461	0.0501	0.9833		1.8
32	0.0420	0.0557	0.9823		1.8
33	0.0402	0.0528	0.9844	***	1.8
34	0.0414	0.0619	0.9833		1.8
35	0.0397	0.0510	0.9792		1.9
36	0.0403	0.0597	0.9781		1.8
37	0.0400	0.0491	0.9854	***	1.8
38	0.0410	0.0704	0.9750		1.8
39	0.0438	0.0802	0.9698		1.8
40	0.0382	0.0532	0.9833		1.9
41	0.0350	0.0517	0.9844		1.8
42	0.0432	0.0516	0.9844		1.8
43	0.0387	0.0561	0.9813		1.8
44	0.0374	0.0660	0.9813		1.8
45	0.0370	0.0540	0.9813		1.8
46	0.0369	0.0544	0.9823		1.8
47	0.0383	0.0509	0.9823		1.8

Training Confusion Matrix =

```
[[ 3616   45]
 [  36 3663]]
```

Training Accuracy = 0.9890

Validation Confusion Matrix =

```
[[ 490    9]
 [  5 456]]
```

Validation Accuracy = 0.9854

Testing Confusion Matrix =

```
[[1027   13]
 [  6 1034]]
```

Testing Accuracy = 0.9909

*** Early stop after 10 epochs with no improvement

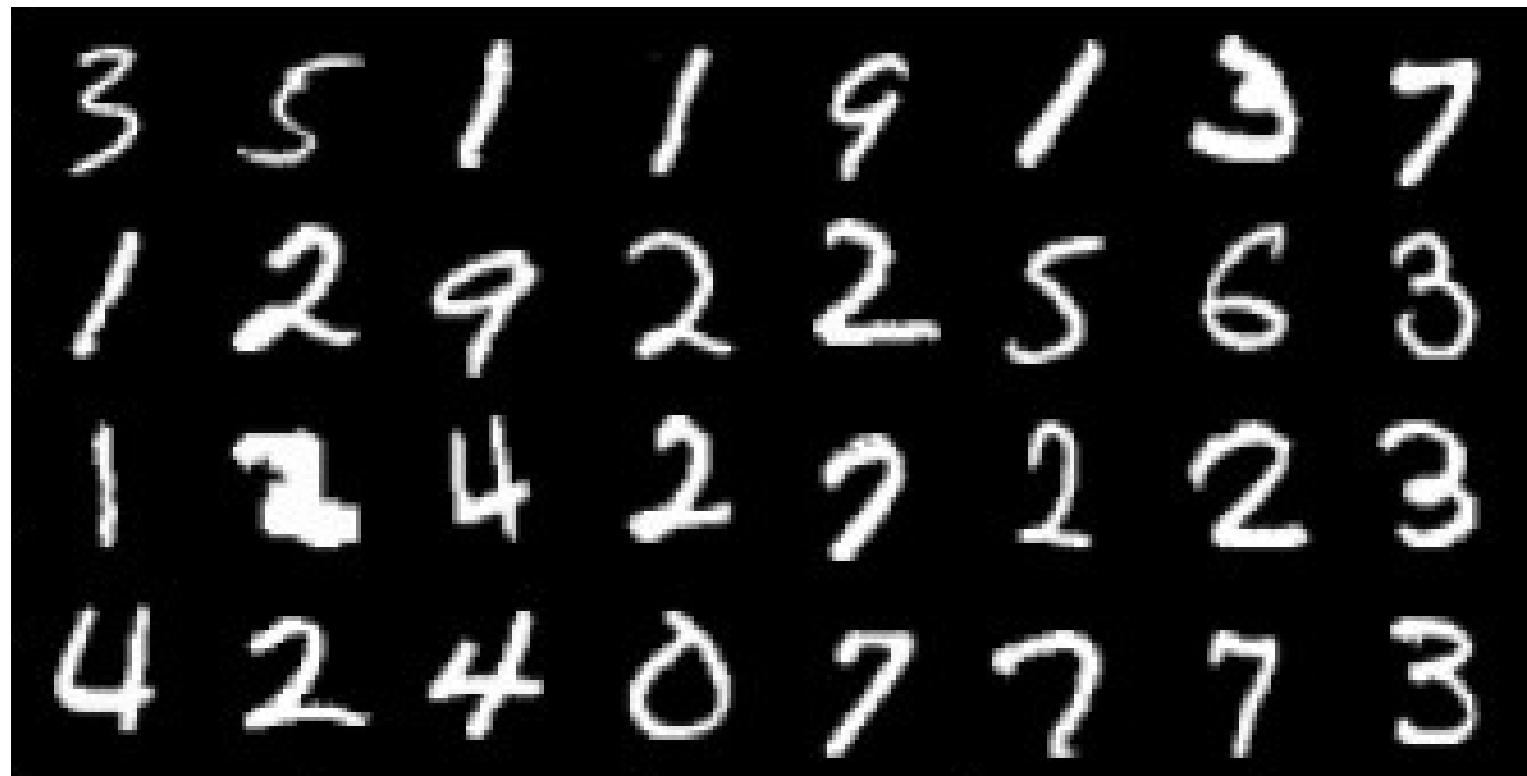
Best model saved best_model.pt (val_acc = 0.9854 in epoch = 37)
Last model saved last_model.pt (val_acc = 0.9823 in epoch = 47)

Fernando Montoya C

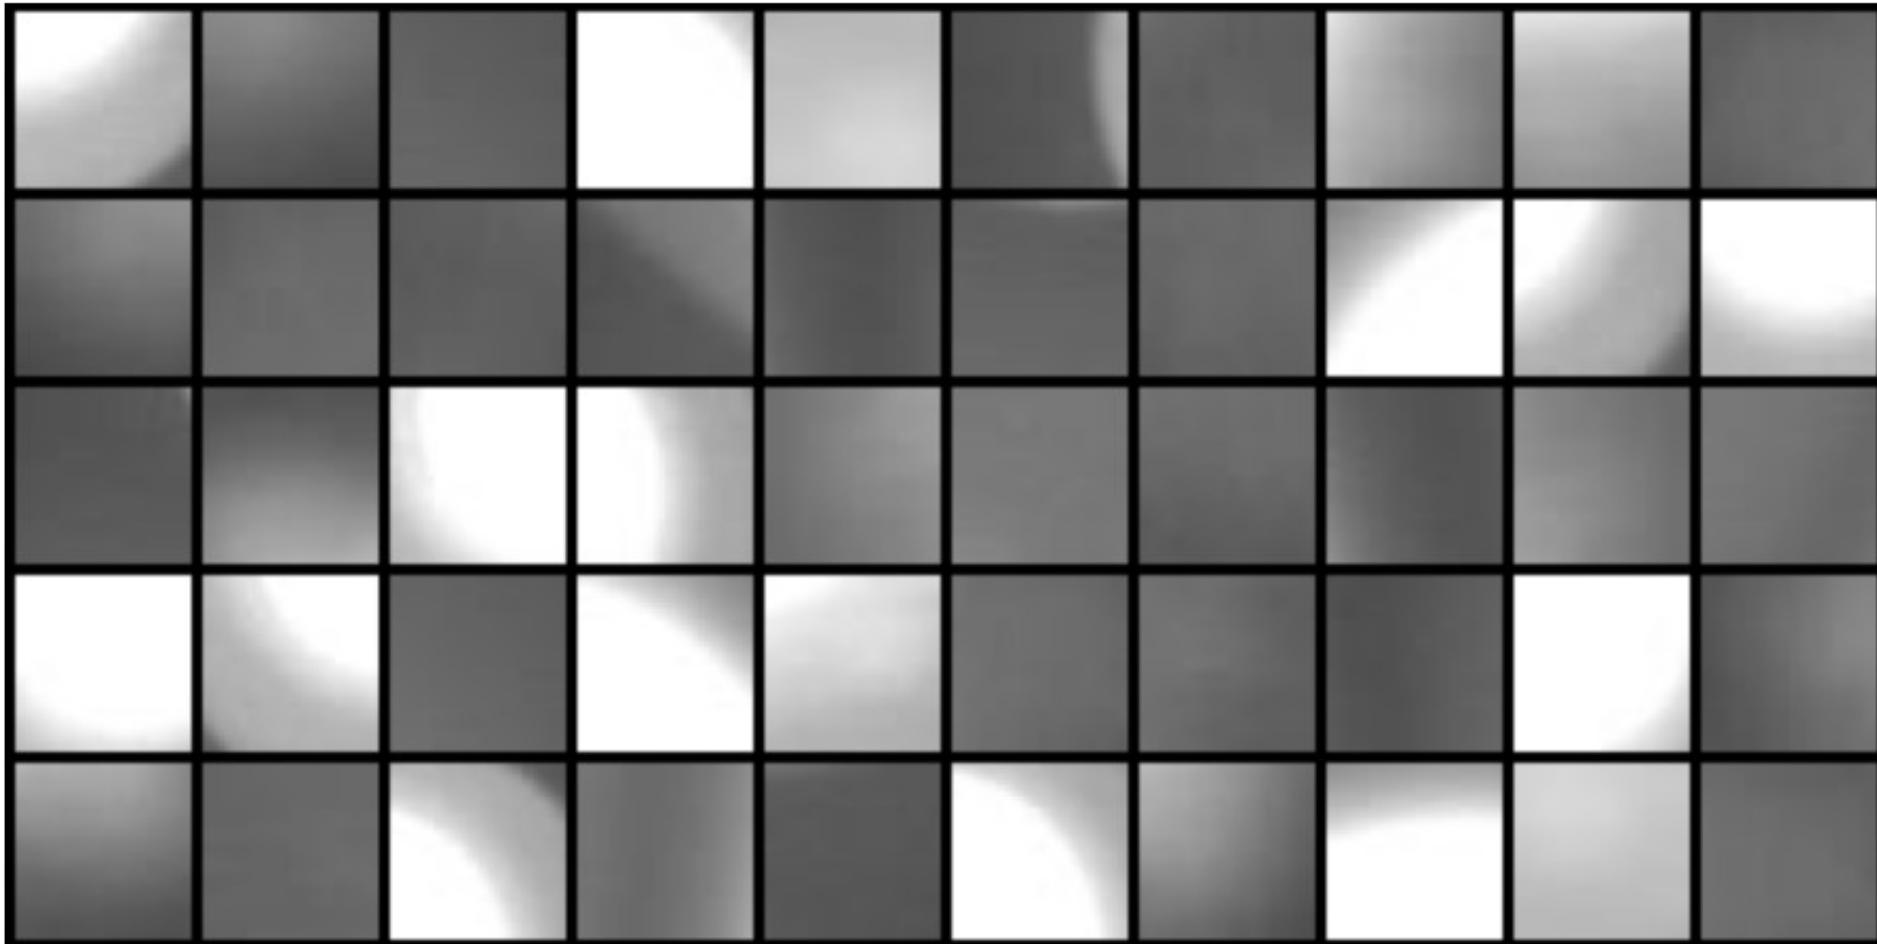
34

Training Time: 85.82 sec

Mas Ejemplos: MNIST



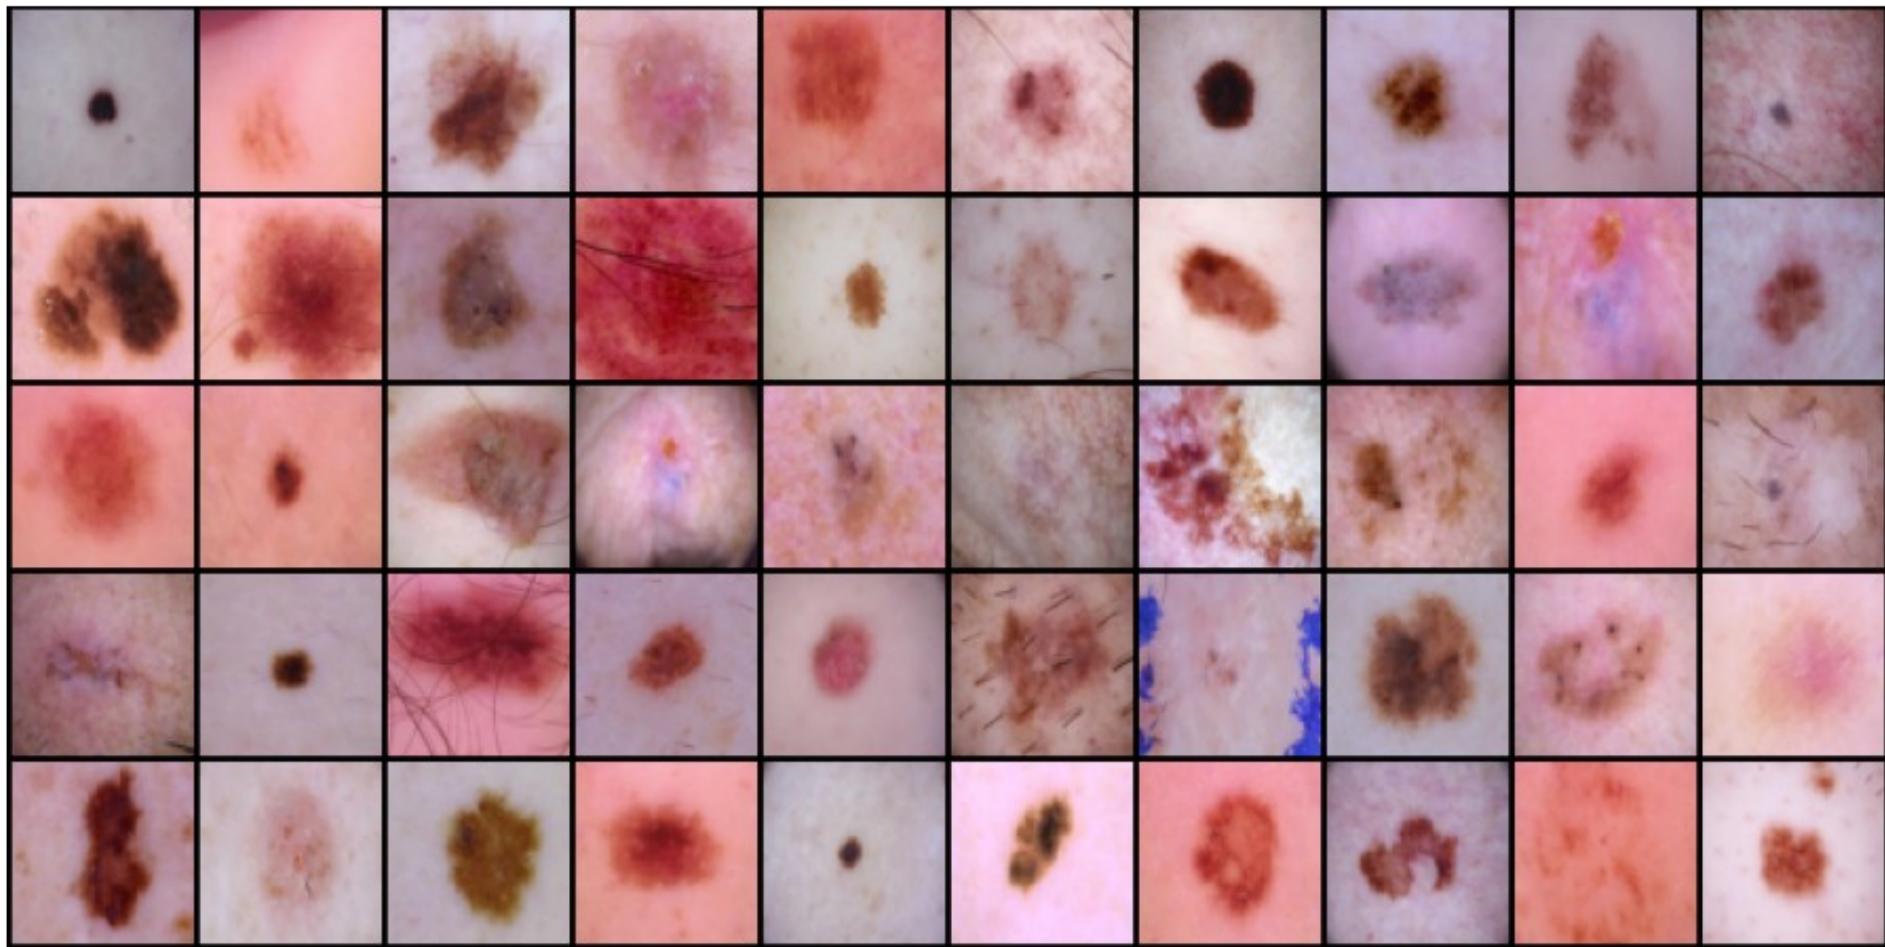
Mas Ejemplos: Defectos



Mas Ejemplos: Gatos vs Perros



Mas Ejemplos: Cáncer de piel 2



Method	Implementation	Features	Gun		Shuriken		Blade		η	Total
			Pr	Re	Pr	Re	Pr	Re		
ResNet50	2	2048	1.00	0.90	1.00	1.00	0.35	0.55	0.80	
VGG16	1	1000	0.95	0.86	0.56	1.00	0.36	0.84	0.73	
VGG19	2	4096	0.78	0.72	0.27	1.00	0.18	0.58	0.53	
AlexNet	2	4096	0.80	0.96	0.84	1.00	0.62	0.54	0.79	
DenseNet121	1	1000	0.99	0.82	0.99	0.98	0.34	0.61	0.78	
GoogleNet	2	1024	0.89	0.98	0.74	1.00	0.23	0.20	0.67	
InceptionResNetV2	0	1536	0.90	0.60	0.82	0.99	0.55	0.60	0.74	
InceptionV3	0	2048	0.85	0.69	0.76	1.00	0.47	0.40	0.69	
MobileNet	2	1280	0.98	0.84	0.94	1.00	0.35	0.95	0.82	
RCNN-ILSVRC13	2	4096	0.75	1.00	0.69	1.00	0.37	0.15	0.64	
ShuffleNet	2	544	0.97	0.78	0.86	1.00	0.20	0.95	0.74	
SqueezeNet	2	1000	0.96	0.38	0.27	1.00	0.06	0.35	0.42	
Xception	0	2048	0.54	0.12	0.75	0.78	0.06	0.21	0.38	
ZfNet512	2	1024	0.87	0.84	0.28	1.00	0.04	0.10	0.48	

GRACIAS

